

## WANG, XU

---

### CONTACT INFORMATION

909, Building New 7, Erqijuchanglu Dongli,  
Xicheng District, Beijing, China  
100045

*Mobile:* +86 13911635527  
*Email:* gnawux@gmail.com  
*Homepage:* <http://wangxu.me/>

### SUMMARY

Backend developer at VisualOps (<http://www.visualops.io> a.k.a MadeiraCloud).

Formerly EBS developer at SNDA Cloud (<http://www.grandcloud.cn>); and developer in Big-Cloud team of China Mobile Research Institute, having nearly 6 year experiences on distributed storage systems such as Hadoop HDFS.

Having been using and maintaining Linux systems for more than 10 years; Developing with C/C++, Scala/Akka, Java, Python, Perl, and Skillful in Unix utilities such as sed.

Technical Blogger, Writer, and Translator on Linux system, NoSQL and Hadoop, etc. Translator of O'Reilly's *Cassandra: The Definitive Guide*.

### CURRENT OCCUPATION

#### Backend Developer in VisualOps

*Backend Developer* (June 2013 – present), currently being responsible for the development and operation of some backend modules:

- *OpsBackend*: Scala/Akka based daemon; communicate with agents on users' instances via Websocket, track state execution, and update recipes when user changed recipe, resource changed, or triggered by API.
- *ResourceDiff*: Scala/Akka based daemon; work with cloud API, discover cloud resources, generate json model of users' VPC/network, and detect changes outside VisualOps.

### PAST PROJECTS AND EXPERIENCE

#### EBS of SNDA GrandCloud

*Senior Researcher (Software Development)* (May 2011 – May 2013), being responsible for the development and operation of EBS (Elastic Block Store) Product. The service was opened to public for testing in late July 2011, and began offering commercial service in Oct 2011. We has been keeping upgrade the stability and performance continuously.

- Jul 2012, *Speech Cassandra: Technical and Historical Overview* in QClub hosted by InfoQ:  
<http://www.slideshare.net/gnawux/cassandratechnicalandhistoryoverview>
- Oct 2011, Hosted the Cloud Computing Track in *QCon Hangzhou 2011*, Hangzhou, China

#### Hadoop Related Work

*Researcher and Project Manager* (July 2007 – May 2011) on Cloud Computing in *China Mobile Research Insititute*, being responsible of *Hadoop HDFS* related development and maintenance, and communicate with upstream community, includes evaluation and testing, bug fix, and feature development (Apache JIRA ID gnawux):

- 2010 – 2011, Feature of DataNode volume refreshment HDFS-1362;
- 2008, Fixed a critical bug of Hadoop 0.18, which may lead to replica mistaken deleting: HADOOP-4742;
- Aprial 2008 – May 2011, Maintaining the CMRI distribution of Hadoop, including NameNodes synchronization patches:  
<http://github.com/gnawux/hadoop-cmri/>

- 2008 – 2009, Maintaining the CMRI Hadoop stress testing tool (Python/Java):  
<http://code.google.com/p/hadoop-test/>
- September 2010, Present in Hadoop in China (HIC 2010), introduced Hadoop related developments in China Mobile Research Institute:  
<http://wangxu.me/blog/p/445>
- November 2009, Present on Apache Roadshow Asia 2009 (Nov, 2009), introduced Hadoop HDFS development in CMRI.
- Weblog on Hadoop related developments: <http://gnawux.info/hadoop/>

### Cloud Storage, Linux & Open Source related

- *Linux Cluster Management*, Since 2010, Having been responsible for system management of the cloud cluster of China Mobile Research Institute, including the configuration maintenance, optimization of virtualization system and Hadoop platform;
- *Cloud Storage*, Since 2010, as Project Manager and Architect, Having been responsible of Cloud Storage Project of China Mobile Research Institute;
- Organizer and Topic Contributor of LinuxFB: Monthly Salon of Linuxer in BUPT.
- L<sup>A</sup>T<sub>E</sub>X Template for Thesis of BUPT, including code, examples, and manual.

### Technical Writer and Translator

- *Translation: Cassandra: The Definitive Guide*, ISBN: 978-7-115-25854-0, Published in Aug 2011;
- *Book: Life in Debian*, ISBN: 978-7-115-19466-4, Published in Sep 2009.
- *Articals*: Techonology Author and Translator on Linux kernel, Filesystem, NoSQL, etc., Articles posted to Personal Web Site (<http://wanxu.me/blog/>) and Magazines.

### Training Lecturer

2005 – 2007

Staff training or certification training on Linux and UMTS/TD-SCDMA Core Networks, Services, and Air interfaces, for operators and vendors, in Beijing, Guangdong, Wuhan, etc.

### Software Developer

2002 - 2003

Having been responsible for development of wireless field test system, etc.

## EDUCATION

### Beijing University of Posts and Telecommunications (BUPT)

Ph.D., Circuits and Systems, 2002.09 — 2007.07

### Beijing University of Technology

B.S., Electronic Engineering, 1998.09 — 2002.07

## REFERENCES

### Part of Articles

Hadoop related technical writing (<http://wangxu.me/blog/p/tag/hadoop>)

- DataNode Volume Refreshment in HDFS-1326 (English):  
<http://gnawux.info/hadoop/2010/12/datanode-volume-refreshment-in-hdfs-1362/>
- How to Decomission DataNodes in HDFS:  
[http://labs.chinamobile.com/mblog/225\\_18378](http://labs.chinamobile.com/mblog/225_18378)
- dfs.datanode.max.xcievers: An Undocumented Parameter of Hadoop:  
[http://labs.chinamobile.com/mblog/225\\_18094](http://labs.chinamobile.com/mblog/225_18094)

Other articles on scripts, frameworks, etc.

- Stack Context of tornado (python web framework by facebook/friendfeed):  
<http://wangxu.me/blog/p/758>
- Update on tornado-memcache: <http://wangxu.me/blog/p/764>
- Words on expect: <http://wangxu.me/blog/p/752>
- Process sub of Bash: <http://wangxu.me/blog/p/738>
- Introduction on ZeroMQ: <http://wangxu.me/blog/p/725>

## Part of Translations

NoSQL, Hadoop, and Java related translations:

- ZooKeeper Watches: <http://wangxu.me/blog/p/546>
- Cassandra by Example: <http://wangxu.me/blog/p/383>
- Observers: Making ZooKeeper Scale Even Further: <http://wangxu.me/blog/p/320>
- Avro: a New Format for Data Interchange: <http://wangxu.me/blog/p/279>
- Java SE 6 Hotspot<sup>TM</sup> Virtual Machine Garbage Collection Tuning :  
<http://wangxu.me/blog/p/209>

Linux kernel, graphics and other software related translations:

- LWN: Flushing out the pdf: <http://wangxu.me/blog/p/610>
- LWN: Improving ext4: bigalloc, inline data, and metadata checksums:  
<http://wangxu.me/blog/p/604>
- LWN: A tale of two SCSI targets: <http://wangxu.me/blog/p/586>
- LWN: The Fun of Null Pointers (1): <http://wangxu.me/blog/p/39>
- KernelTrap: Linux: The Journaling Block Device: <http://wangxu.me/blog/p/151>
- LWN: What is RCU, Fundamentally?: <http://wangxu.me/blog/p/122>

Note: All the original articles are from Internet, and announced the info about original articles; and most articles, including those from LWN, cloudera has asked the authorization of original publishers